

Supplemental Material

Additional performance analysis results

In Table 1, we list the timing for all three evaluated steps (linearization and shader recompilation, visibility mask computation, and rendering) for all graphics cards on all the datasets. On the Nvidia Quadro RTX 8000, the DOS method still runs at around 2 frames per second (FPS) at a 2048×2048 resolution with 1024 slices, and at 4 to 7 FPS at a 1024×1024 resolution with 512 slices even for the largest datasets (fibers and mitochondria).

	Frame size	Number of slices		
		256	512	1024
Integrated graphics				
Fibers	512×512	6 / 468 / 429	5 / 468 / 786	6 / 468 / 1502
Pores	1024×1024	6 / 469 / 1012	7 / 468 / 1917	6 / 468 / 3722
Mitosis	2048×2048	9 / 471 / 2736	10 / 471 / 5210	9 / 472 / 10170
Fibers	512×512	5 / 545 / 464	6 / 556 / 839	5 / 545 / 1629
Pores	1024×1024	6 / 533 / 934	7 / 535 / 1760	6 / 542 / 3439
Mitosis	2048×2048	9 / 551 / 2618	8 / 551 / 4990	9 / 545 / 9772
Fibers	512×512	6 / 546 / 487	6 / 545 / 898	5 / 546 / 1734
Pores	1024×1024	5 / 547 / 985	7 / 546 / 1839	6 / 545 / 3655
Mitosis	2048×2048	8 / 550 / 2642	8 / 552 / 5036	8 / 548 / 9825
Commodity desktop graphics				
Fibers	512×512	9 / 89 / 221	11 / 89 / 421	15 / 89 / 805
Pores	1024×1024	12 / 89 / 215	15 / 89 / 373	12 / 89 / 689
Mitosis	2048×2048	14 / 89 / 405	15 / 89 / 617	13 / 89 / 1203
Fibers	512×512	7 / 56 / 161	15 / 47 / 287	18 / 52 / 542
Pores	1024×1024	21 / 48 / 172	12 / 52 / 294	19 / 50 / 523
Mitosis	2048×2048	13 / 48 / 356	15 / 48 / 597	15 / 49 / 1068
Fibers	512×512	9 / 26 / 101	9 / 29 / 187	11 / 27 / 327
Pores	1024×1024	9 / 27 / 134	12 / 27 / 215	15 / 27 / 383
Mitosis	2048×2048	17 / 27 / 340	17 / 27 / 575	18 / 29 / 1018
Professional graphics				
Fibers	512×512	5 / 85 / 34	7 / 85 / 55	5 / 85 / 100
Pores	1024×1024	5 / 85 / 81	5 / 85 / 138	4 / 85 / 248
Mitosis	2048×2048	4 / 86 / 204	4 / 87 / 324	4 / 91 / 580
Fibers	512×512	5 / 47 / 26	5 / 47 / 41	4 / 47 / 72
Pores	1024×1024	4 / 47 / 60	4 / 46 / 92	4 / 46 / 156
Mitosis	2048×2048	4 / 46 / 162	4 / 46 / 231	4 / 50 / 368
Fibers	512×512	5 / 26 / 39	6 / 26 / 67	5 / 25 / 123
Pores	1024×1024	4 / 25 / 91	4 / 25 / 160	4 / 25 / 294
Mitosis	2048×2048	4 / 25 / 203	4 / 25 / 326	4 / 25 / 574

Table 1: Performance evaluation of the volume conductor. The three numbers are the times in milliseconds to linearize the predicates and recompile the shader, recompute the visibility mask, and render the image.

We also present additional graphs for integrated graphics (Fig. 1) and commodity desktop graphics (Fig. 2).

Algorithms

Algorithms Algorithm 1 and Algorithm 2 describe the instance sparsification and visibility mask generation processes in detail.

Online repository and examples

The source code and all examples with synthetic data are available online at <http://lgm.fri.uni-lj.si/portfolio-view/volume-conductor>. The online demo is implemented in WebGL 2.0 Compute. While the API is unsupported in stable releases of web browsers, the demo can be tested in

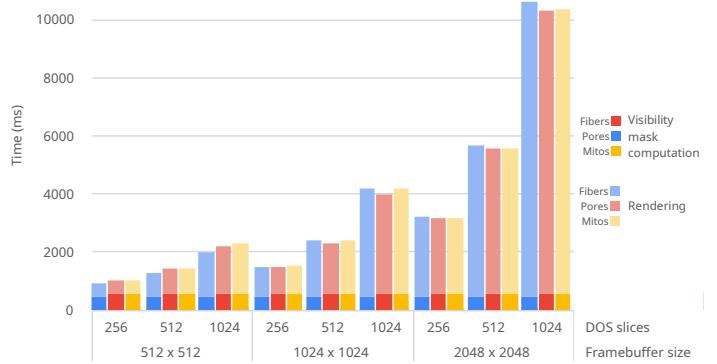


Fig. 1: Performance evaluation on a laptop computer with integrated graphics. Linearization and shader recompilation times are negligible and not depicted in the graph.

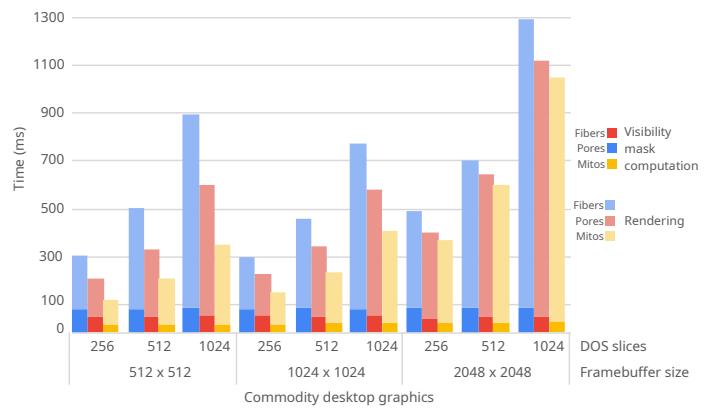


Fig. 2: Performance evaluation on a commodity desktop computer with an Nvidia GTX 1060 graphics card. Linearization and shader recompilation times are negligible and not depicted in the graph.

Chromium Dev build 87 available at <https://commondatastorage.googleapis.com/chromium-browser-snapshots/index.html>.

Additional examples

This section contains additional examples of the features of the volume conductor. Fig. 3 presents additional renders of the fibers, pores, and mitochondria. Fig. 4 shows the effects of different sparsification functions on a synthetic dataset. In Fig. 5 and Fig. 6, we used another synthetic dataset containing boxes, spheres, and ellipsoids of various sizes. The distribution of most instances is uniform, except for ellipsoids, which are placed along one of the spatial diagonals of the volume. A large sphere is placed in the center of the volume, which is not visible without sparsification. The volume conductor is designed to be independent of the rendering method; thus, we include a comparison of path tracing and directional occlusion shading (Fig. 7).

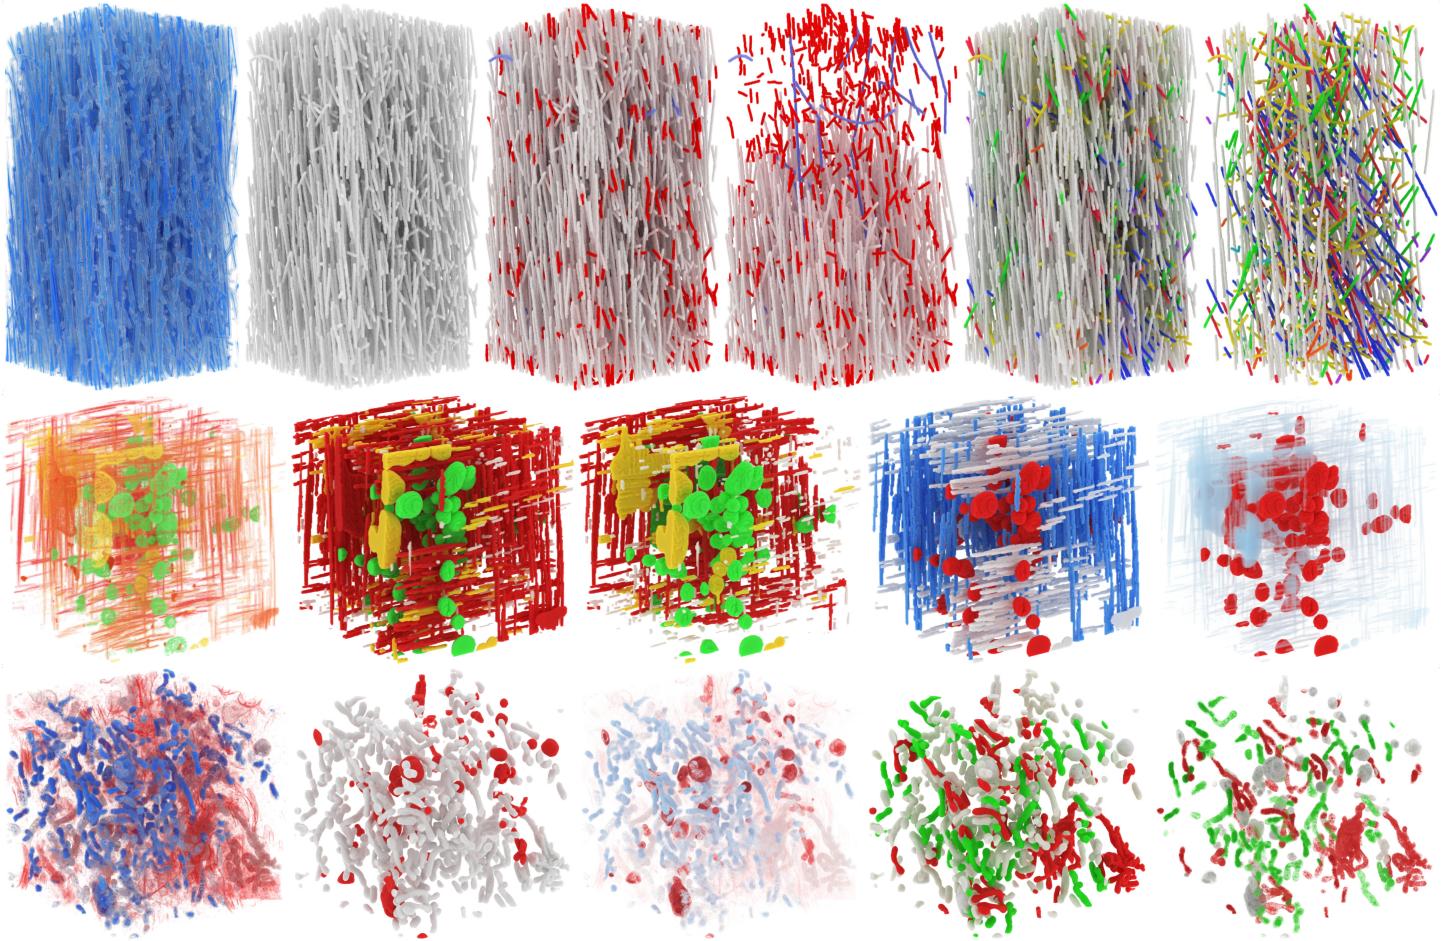


Fig. 3: Top row, fibers: raw data, segmentation, colorization of short (red) and bent (blue) instances, same as previous but with 50 % of the remaining instances hidden, colorization by orientation, the same as previous with 80 % of the vertical fibers hidden. **Middle row, pores:** raw data, colorization by roundedness, the same as previous with 50 % of the needle-shaped pores hidden, colorization of needle-shaped pores by orientation, and ghosting. **Bottom row, mitochondria:** raw data, colorization of type, the same as previous but with added blending with raw data, colorization by branching and thinning, same as previous with opacity transfer. Path tracing was used to generate the images.

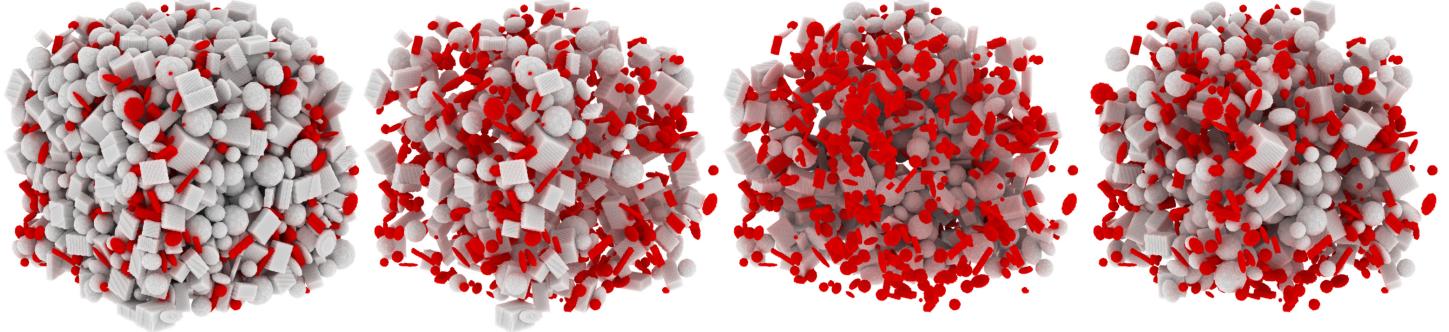


Fig. 4: Demonstration of sparsification functions. From left to right: no sparsification, random, depth-based, and context-preserving sparsification. Instances with a small volume are colored red. Path tracing was used to generate the images.

Algorithm 1 Sparsification with temporal coherency.

```

1: for all  $i \in \text{instances}$  do
2:    $hidden_i \leftarrow \text{false}$ 
3: end for
4: for all  $r \in \text{predicates}$  do
5:    $G \leftarrow \text{instances}(r)$ 
6:    $n_{\text{hidden}} \leftarrow 0$ 
7:    $n_{\text{toHide}} \leftarrow \lfloor \text{visibilityRatio}(r) \cdot |G| \rfloor$ 
8:    $\text{sort}(G)$ 
9:   for all  $i \in G$  do
10:    if not  $hidden_i$  and not  $visible_i$  then
11:       $n_{\text{hidden}} \leftarrow n_{\text{hidden}} + 1$ 
12:       $hidden_i \leftarrow \text{true}$ 
13:      if  $n_{\text{hidden}} = n_{\text{toHide}}$  then
14:        break
15:      end if
16:    end if
17:  end for
18:  for all  $i \in G$  do
19:    if  $hidden_i$  then
20:      continue
21:    else if  $n_{\text{hidden}} \leq n_{\text{toHide}}$  then
22:       $n_{\text{hidden}} \leftarrow n_{\text{hidden}} + 1$ 
23:       $visible_i \leftarrow \text{false}$ 
24:    else
25:       $visible_i \leftarrow \text{true}$ 
26:    end if
27:  end for
28: end for

```

Algorithm 2 Visibility mask generation.

```

1: for all  $i \in \text{instances}$  do
2:   for all  $r \in \text{predicates}$  do
3:     if  $r(i)$  then
4:       if  $visible_i$  then
5:          $group_i \leftarrow r$ 
6:          $maskValue_i \leftarrow \text{maskValue}(r)$ 
7:       else
8:          $group_i \leftarrow 0$ 
9:          $maskValue_i \leftarrow \text{maskValue}(0)$ 
10:      end if
11:    end if
12:  end for
13: end for
14: for all  $v \in \text{voxels}$  do
15:    $i \leftarrow \text{id}(v)$ 
16:    $maskValue_v \leftarrow maskValue_i$ 
17: end for

```

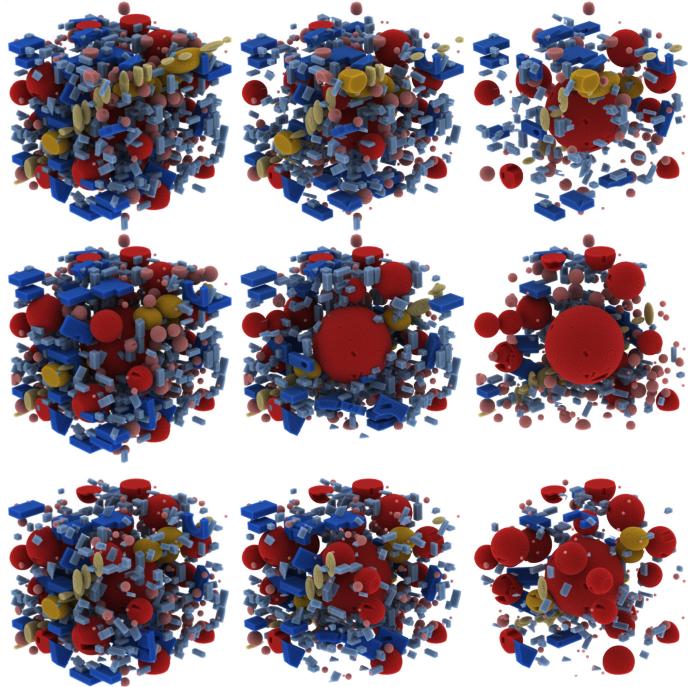


Fig. 5: From left to right: increasing degrees of uniform (top row), depth-based (middle row), and context-preserving (bottom row) sparsification.

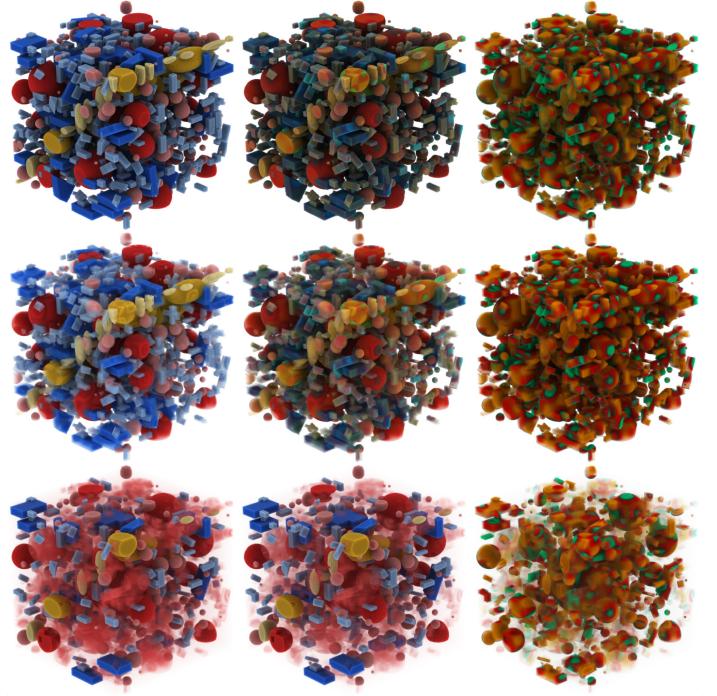


Fig. 6: Top row: different colorization as a result of blending the raw data and segmentation data transfer functions. Middle row: the same colors as above, but with the opacity transferred from the raw data transfer function, enabling the user to employ both the sparsification features of the volume conductor and the opacity set by the transfer function. Bottom row: ghosting of the instances as a result of blending the opacity after sparsification.

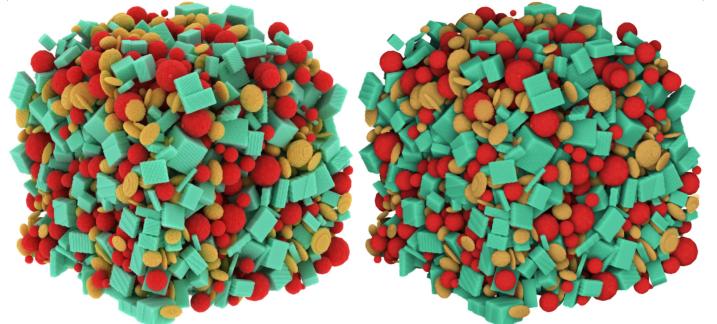


Fig. 7: Left: Rendering comparison of path tracing; Right: directional occlusion shading.